

MyMentalHealthBuddy Release handoff

September 15, 2026 | Scope: MyMentalHealthBuddy.com only

The website is live. Full production readiness remains unverified. This review identified the serving Replit app, checked public behavior, prepared one bounded correction, and preserved prior work. It did not deploy a new version.

Area	Evidence and status
Published site	Replit reports a successful deployment at www.mymentalhealthbuddy.com . Public home and crisis pages rendered.
Runtime checks	<code>/health</code> , <code>/ready</code> and <code>/api/health</code> returned HTTP 200 JSON. These responses do not establish database readiness or private feature correctness.
NAMI directory	Live page shows the old NAMI text destination. Correction prepared; 13 component checks and 13 repair-script checks passed.
Build	Baseline and candidate builds passed. The current Replit workspace has not been built in this review.
Library navigation	FAIL: the Health learning link opens <code>/health</code> diagnostic JSON. This is the next queued repair.
Private features	Login, account isolation, journal/mood persistence, payments and recovery remain unverified in this review.
Content and assistant	26-category index, 3 complete draft lessons, 30-day editorial CSV and 8-slide deck preserved. Existing daily MMHB draft task is enabled.
Publishing changes	GitHub push blocked by automatic approval review. No new CI run or production deployment completed.

A completion percentage would be misleading: an untested private workflow can still be a release blocker even when most pages exist. Use the acceptance checks on page 3 to measure readiness.

The prepared correction

The NAMI HelpLine entry now points to Text NAMI to 62640. It describes non-crisis emotional support, links directly to the official HelpLine page, and includes Monday-Friday hours, 10am-10pm Eastern, excluding federal holidays. [1]

The correction changes four strings in client/src/pages/CrisisResources.jsx. All six support cards, voice actions, 988 guidance, and the separate Crisis Text Line instruction remain intact. The repair adds no dependencies and does not change routes, styles, authentication or data storage.

Your next Replit action

1. Open your existing Replit app using the link below. The app title is TheGenuineLoveProject; its confirmed published domain is MyMentalHealthBuddy.com. This title does not change the project scope.

Existing Replit app: <https://replit.com/@MyMentalHealth/TheGenuineLoveProject>

2. Upload mmhb-nami-repair.mjs to the project root, beside package.json. Open Shell in that same project root. Run:

```
node ./mmhb-nami-repair.mjs --inspect
```

3. Return the complete JSON output. It contains selected source hashes, Git metadata and the repair eligibility. It does not read .env files, print credentials, contact services, or alter application files.

Reported eligibility	Meaning
EXACT_OLD_BLOCK_MATCH	The known NAMI record is present and the prepared repair fits.
ALREADY_CORRECT	The corrected record is already present. No repeat repair is needed.
SOURCE_DIFF_REVIEW_REQUIRED	The record differs. Preserve the file and adapt the patch to its actual contents.

The included --apply mode is available for the reviewed next step. It changes only the matching NAMI record and verifies a backup first. The output names its backup folder. --rollback BACKUP_FOLDER restores only if the file still matches the recorded repaired version. Avoid concurrent editing. This script does not build, commit, push or deploy.

Release checks that remain

Complete these against the current Replit source and a candidate built from that source. Keep the tested source commit, artifact identity and deployed version together.

Gate	Required evidence
Current source	Current Replit branch, HEAD, tracked modifications and relevant source hashes; reconcile against GitHub without overwriting newer edits.
Candidate build	Use the existing lockfile and actual build script. Build outside the currently served output folder. Confirm the exact changed files.
Account boundaries	Synthetic users A and B: login/logout/reset/expired session, anonymous denial, wrong-role denial, and denied access to the other user's objects.
Saved records	Save, reopen, edit, export and delete synthetic journal/mood records. Restart the candidate and confirm persistence.
Crisis and AI	Verify each exposed crisis contact and region. Test urgent-support visibility, provider outages, unsupported medical claims, prompt injection, privacy, rate limits and spending controls.
Navigation and access	Repair the Health learning link without disrupting deployment health checks. Verify direct URL refresh, keyboard use, focus, mobile menu and text reflow.
Payments, if exposed	Use test mode to prove webhook signatures, repeated-event handling, entitlements and failure recovery.
Recovery	Restore a backup into an isolated database. Rehearse rollback and define acceptable data loss and recovery time.
Release	Required CI passes on the exact candidate. Deploy that artifact, confirm current crisis copy and user journeys, and observe errors before calling it stable.

Known next defect: the visible Wellness learning library link labeled Health navigates to /health and displays runtime JSON. Its replacement route must be selected from the current route registry. Do not repurpose or remove a deployment health endpoint blindly.

Content you can teach

The preserved starter package has three complete original draft lessons: small-step, facts-and-interpretations, and attention-and-choice. The A-Z index has 26 categories; its remaining entries are explicitly marked as requiring research. It is not an exhaustive or fully reviewed clinical library.

Teaching level	Example: a goal and a next step
1. Understand	A goal names a direction: I want an organized desk. A next step names something I can do: put three papers in a folder.
2. Try	After breakfast, at my desk, I will put three papers in a folder. If that is too much, choose a smaller action or ask for help.
3. Teach back	Ask: How would you explain the difference? Example answer: A goal names the direction; a next step names an action, cue and place.

Original metaphor: a large goal can resemble a staircase, while a next step gives one place to focus. Its limit: real life may call for a ramp, rest, support or another route. This is an illustration, not a neurological explanation or treatment claim.

The editable eight-slide deck includes concrete comparisons, an optional exercise and speaker notes. Its general self-care and mindfulness sources were reopened in this review. Those sources do not establish the effectiveness of these exact scripts. [2,3]

Daily workflow

Use the existing MMHB draft assistant around 6pm America/Los_Angeles. Spend five minutes choosing one topic, ten drafting one example, ten adapting the format and five reviewing sources. Review each draft, post through the native social app, then record its actual URL and metrics in your editorial plan. The task produces conversation drafts; it does not publish posts.

A future website calendar should save the topic, owner, version, source checks, reviewer, intended timezone, destination and status. Separate draft, approved, scheduled, publishing, published, failed and canceled states. Confirm a provider post ID before marking an item published. This is a design requirement, not an implemented feature claim.

Support choice through open questions and accurate reflections. Keep philosophy and named spiritual traditions clearly identified. NLP language exercises, shadow-work metaphors and quantum terminology must not be presented as proven treatments or promises of healing. Keep affiliate links separate from private reflection and crisis support.

Three stages of completion

Stage	Exit condition	Provisional effort
1. Dependable release	Current source reconciled; required private and public journeys, recovery and deployment gates pass.	24-60 engineer-hours
2. Creator workspace	Initial reviewed library; durable editor and calendar; one provider; generation, review, scheduling and duplicate-post protection tested.	Additional 60-120 hours
3. Measured expansion	Broader research and media, verified affiliate catalog and analytics; one reversible experiment at a time.	Additional 80-180 hours

Combined planning range: 164-360 engineer-hours for these finite stages. These are estimates, not measured remaining work or delivery dates. They assume the existing architecture is reusable and required accounts are available. They exclude external approval and clinical review wait times; deeper defects may exceed the range. Re-estimate after current-source and two-account tests.

Integration responsibilities

Service	Proposed role and current limit
OpenAI / Codex	Use Codex for bounded code changes and evaluation. Use the server-side Responses web_search capability for cited research if selected. No new website API integration was executed. [4]
Perplexity	Optional second research provider. Its Sonar API supports grounded responses. No app keys or paid calls were configured. [5]
PubMed / DailyMed	PubMed for literature discovery; DailyMed for label references. Preserve identifiers, versions and source dates. No production ingestion completed. [6,7]
CMS	Separate coverage decisions and public administrative datasets from clinical evidence. No coverage or CMS data integration was completed. [8,9]
Manual social workflow	Metricool will not be used for this request. Review drafts and publish through native social apps. Record actual post URLs and exported metrics; current trends cannot be ranked without account data.
Books / affiliates	Use licensed assets and original summaries. Verify program membership and product links; display commission disclosures near links. No affiliate configuration was completed.

A connected ChatGPT plugin and a working website integration are separate achievements. Website integration still needs application credentials, access checks, durable storage, errors, limits and end-to-end tests.

Reusable execution prompts

Build continuation: Work only on MyMentalHealthBuddy. Use the attached current-source report, not earlier assumptions. Identify one observable failing user journey. Prove the responsible route or file, prepare the smallest patch and preserve current work. Use shell commands and existing dependencies. Compare baseline and candidate behavior, record the exact diff, test the relevant account and data boundaries, and prepare rollback. Report local verification, CI, deployment and production evidence separately. Do not mark unknown gates passed.

Lesson generation: Create one original adult education lesson from sources actually read. State the audience, learning objective, evidence category and limits. Explain the idea plainly; add an everyday example, an explicitly labeled metaphor and its limit, one optional practice, an open question and a teach-back answer. Produce platform-specific draft variants, a visual brief and alt text. Keep citations near claims. Do not diagnose, promise cures, infer trauma histories, fabricate trends, copy books or invent affiliate links. Return DRAFT and identify what needs review.

Trend review: Using only supplied or retrieved account metrics, compare posts on the same platform and format at comparable ages. Report sample sizes, missing data, completion, saves, shares, site visits and negative feedback. Recommend one reversible test with a hypothesis and primary measure. Do not infer causation from one successful post. If account data is absent, return DATA_UNAVAILABLE rather than inventing a ranking.

References and verification limits

[1] NAMI HelpLine: <https://www.nami.org/nami-helpline/>

[2] NIMH, Caring for Your Mental Health: <https://www.nimh.nih.gov/health/topics/caring-for-your-mental-health>

[3] NCCIH, Meditation and Mindfulness: <https://www.nccih.nih.gov/health/meditation-and-mindfulness-effectiveness-and-safety>

[4] OpenAI, Web search: <https://developers.openai.com/api/docs/guides/tools-web-search>

[5] Perplexity, Sonar API: <https://docs.perplexity.ai/docs/sonar/quickstart>

[6] NCBI, E-utilities introduction: <https://www.ncbi.nlm.nih.gov/books/NBK25497/>

[7] DailyMed, Web services: <https://dailymed.nlm.nih.gov/dailymed/app-support-web-services.cfm>

[8] CMS, Coverage API: <https://api.coverage.cms.gov/>

[9] CMS, Public data API: <https://data.cms.gov/api-docs>

Reference review: September 15, 2026; NCBI and CMS public-data roles were confirmed from primary-source search results. Local candidate visual inspection was blocked because the browser could not open localhost. The component render stubs surrounding wrappers. Repository TypeScript checks do not fully cover JSX. No full clinical review, penetration test, restoration exercise, current social analytics study or production deployment was performed.